

Car Price Prediction

August 26, 2025

```
[30]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Lasso
from sklearn import metrics
```

```
[2]: car_data = pd.read_csv('car_data.csv')
```

#First 5 rows of the dataframe

```
[4]: car_data.head()
```

```
[4]:   Car_Name  Year  Selling_Price  Present_Price  Kms_Driven  Fuel_Type  Seller_Type
Transmission  Owner
0    ritz    2014         3.35         5.59        27000    Petrol    Dealer
Manual      0
1    sx4    2013         4.75         9.54        43000    Diesel    Dealer
Manual      0
2    ciaz    2017         7.25         9.85         6900    Petrol    Dealer
Manual      0
3  wagon r    2011         2.85         4.15         5200    Petrol    Dealer
Manual      0
4  swift    2014         4.60         6.87        42450    Diesel    Dealer
Manual      0
```

#Checking the number of rows and columns

```
[5]: car_data.shape
```

```
[5]: (301, 9)
```

```
[6]: #getting some information about the dataframe
car_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype

```

```

---  -----  -----  -----
0   Car_Name      301 non-null  object
1   Year          301 non-null  int64
2   Selling_Price 301 non-null  float64
3   Present_Price 301 non-null  float64
4   Kms_Driven    301 non-null  int64
5   Fuel_Type     301 non-null  object
6   Seller_Type   301 non-null  object
7   Transmission  301 non-null  object
8   Owner         301 non-null  int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB

```

```

[7]: #Checking Missing Value
car_data.isnull().sum()

```

```

[7]: Car_Name      0
Year            0
Selling_Price   0
Present_Price   0
Kms_Driven      0
Fuel_Type       0
Seller_Type     0
Transmission    0
Owner           0
dtype: int64

```

```

[9]: #Checking the distribution of categorical data
print(car_data.Fuel_Type.value_counts())
print(car_data.Seller_Type.value_counts())
print (car_data.Transmission.value_counts())

```

```

Fuel_Type
Petrol    239
Diesel    60
CNG        2
Name: count, dtype: int64
Seller_Type
Dealer    195
Individual 106
Name: count, dtype: int64
Transmission
Manual    261
Automatic  40
Name: count, dtype: int64
#Encoding the category Data

```

```
[14]: #encoding "Fuel_Type" Column
car_data.replace({'Fuel_Type':{'Petrol':0,'Diesel':1,'CNG':2}},inplace=True)

#encoding "Seller_Type" Column
car_data.replace({'Seller_Type':{'Dealer':0,'Individual':1}},inplace=True)

#encoding "Transmission" Column
car_data.replace({'Transmission':{'manual':0,'Automatic':1}},inplace=True)
```

```
[15]: car_data.head()
```

```
[15]: Car_Name  Year  Selling_Price  Present_Price  Kms_Driven  Fuel_Type
Seller_Type  Transmission  Owner
0    ritz    2014           3.35           5.59        27000           0
0           0           0
1    sx4    2013           4.75           9.54        43000           1
0           0           0
2    ciaz    2017           7.25           9.85         6900           0
0           0           0
3  wagon r    2011           2.85           4.15         5200           0
0           0           0
4    swift    2014           4.60           6.87        42450           1
0           0           0
```

#splitting the data

```
[19]: x = car_data[['Year', 'Present_Price', 'Kms_Driven', 'Fuel_Type',
↳ 'Seller_Type', 'Transmission', 'Owner']]
y = car_data['Selling_Price']
```

```
[20]: print(x)
```

```
Year  Present_Price  Kms_Driven  Fuel_Type  Seller_Type  Transmission
Owner
0    2014           5.59        27000           0           0           0
0
1    2013           9.54        43000           1           0           0
0
2    2017           9.85         6900           0           0           0
0
3    2011           4.15         5200           0           0           0
0
4    2014           6.87        42450           1           0           0
0
..    ...           ...           ...           ...           ...
...
296  2016          11.60       33988           1           0           0
0
```

297	2015	5.90	60000	0	0	0
0						
298	2009	11.00	87934	0	0	0
0						
299	2017	12.50	9000	1	0	0
0						
300	2016	5.90	5464	0	0	0
0						

[301 rows x 7 columns]

```
[21]: print(y)
```

```
0      3.35
1      4.75
2      7.25
3      2.85
4      4.60
```

```
...
296     9.50
297     4.00
298     3.35
299    11.50
300     5.30
```

Name: Selling_Price, Length: 301, dtype: float64

#Splitting Training and Test Data

```
[22]: x_train,x_test, y_train, y_test = train_test_split(x, y, test_size =0.1,
↳random_state=2)
```

#Model_Training

#Loading the linear regression model

```
[33]: lin_reg_model = LinearRegression()
```

```
[37]: lin_reg_model.fit(x_train, y_train)
```

```
[37]: LinearRegression()
```

#Model Evaluation

```
[39]: #prediction on Training data
training_data_prediction = lin_reg_model.predict(x_train)
```


#Rsquared error

```
[41]: error_score = metrics.r2_score(y_train, training_data_prediction)
print("Rsquared error : ", error_score)
```

Rsquared error : 0.8799451660493708

#Visualize the actual and Predicted prices:

```
[42]: plt.scatter(y_train, training_data_prediction)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual Prices Vs Predicted Prices")
plt.show()
```




```
[43]: #prediction on Training data
test_data_prediction = lin_reg_model.predict(x_test)
```

```
[45]: error_score = metrics.r2_score(y_test, test_data_prediction)
print("Rsquared error : ", error_score)
```

Rsquared error : 0.8365766715026374

```
[46]: plt.scatter(y_test, test_data_prediction)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
```

```
plt.title("Actual Prices Vs Predicted Prices")
plt.show()
```



#Lasso Regression

```
[47]: lass_reg_model = Lasso()
```

```
[48]: lass_reg_model.fit(x_train, y_train)
```

```
[48]: Lasso()
```


```
[49]: #prediction on Training data
training_data_prediction = lass_reg_model.predict(x_train)
```

```
[50]: error_score = metrics.r2_score(y_train, training_data_prediction)
print("Rsquared error : ", error_score)
```

Rsquared error : 0.8427856123435794

```
[51]: plt.scatter(y_train, training_data_prediction)
plt.xlabel("Actual Price")
```

```
plt.ylabel("Predicted Price")
plt.title("Actual Prices Vs Predicted Prices")
plt.show()
```



```
[52]: #prediction on Training data
test_data_prediction = lass_reg_model.predict(x_test)
```

```
[53]: error_score = metrics.r2_score(y_test, test_data_prediction)
print("Rsquared error : ", error_score)
```

Rsquared error : 0.8709167941173195

```
[54]: plt.scatter(y_test, test_data_prediction)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual Prices Vs Predicted Prices")
plt.show()
```



[]: